

Prof. Stefan Roth
Junhwa Hur
Anne Wannenwetsch

This assignment is due on February 11th, 2019 at 13:00.

Common function module

For this assignment please make use of the functions given in `Common.jl`, which is imported at the top of each Julia problem file.

Packages

You may need to install the Julia packages `Clustering`, `MultivariateStats`, and `Optim`.

Please refer to the previous assignments for general instructions and follow the handin process described there.

In this assignment you will implement a simple image classifier that categorizes images as either airplanes or motorbikes. To classifying images, we rely on two basic design choices:

1. Which *feature representation* do we use to encode our images?
2. Which *classification model* do we use to classify the encoded images?

Figure 1 below shows the process of classifying an image. In this assignment we will use the Bag-of-Words (BoW) representation for describing an image and a Multi Layer Perceptron (MLP) for the classification part.




Figure 1: An image is first turned into a feature vector that is then subsequently classified.

Problem 1 - Bag-of-Words Model (15 points)

In this problem you will implement a simplified Bag-of-Words model. The main idea of the BoW representation is to first detect image features on the image and then use them to encode the image as a distribution over features of a codeword dictionary. Therefore, the image features are each assigned to the closest features of the dictionary and a histogram is used to represent the distribution. With this strategy, we can turn an arbitrary number of image features into a fixed-size histogram. This suits subsequent classification as most classifiers assume a fixed-size input.

Creating the codeword dictionary:

First, we create the codeword dictionary. It is a set of k key point descriptors and we obtain it by applying k -means clustering on *all* feature vectors from *all* the training images. K -means clustering aims to partition n observations $(\mathbf{x}_1, \dots, \mathbf{x}_n)$ into k ($k \leq n$) clusters (s_1, \dots, s_k) as to minimize the within-cluster sum of squares

$$\sum_{i=1}^k \sum_{\mathbf{x}_j \in s_i} \|\mathbf{x}_j - \boldsymbol{\mu}_i\|^2, \quad (1)$$

where $\boldsymbol{\mu}_i$ is the mean of points in s_i .

Tasks:

- You are given image folders for both airplanes and motorbikes. Please implement `loadimages` that creates both a training and a testing structure containing the images, class labels, and number of samples per set. You should randomly separate the two sets of images (*i.e.* including airplanes or motorbikes, respectively) into two equal-size parts for training and testing. Please denote motorbikes with class label $y = 0$ and airplanes with class label $y = 1$.
- Apply the Harris detector and SIFT descriptor to obtain feature vectors for all images in both training and testing set. Use the parameters $\sigma = 1.4$, filter size 15×15 and threshold 10^{-7} for interest point detection and ignore points within a 10-pixel wide boundary. Please implement `extractfeatures` that finds interest points and extracts SIFT descriptors at the detected interest points for each image. You can use the Harris corner detector, `detect_interestpoints`, and the SIFT descriptor, `sift`, implemented in “Common.jl”.
- To get a database of image features, concatenate the features of all images in the training set by implementing the function `concatenatefeatures`. The resulting feature matrix will be used to compute the codebook of key points.
- Implement `compute_codebook` that computes the codeword dictionary using k -means clustering with $k = 50$. The entries of the codebook are given as the centers of the k clusters. You can use the package “Clustering.jl” that already implements k -means.
- For all images you need to compute the histogram over assigned codewords. Therefore, implement the function `compute_histogram` that takes a set of images as well as a codebook and computes a histogram for each image. Here, the histogram displays the number of occurrences of the codewords in an individual image. Use the (squared) Euclidean distance to assign feature vectors to the closest codeword in the codebook and normalize the histogram in the end such that it sums to one. Finally, the function `compute_histogram` returns a so-called feature matrix in which each column contains the histogram of one of the images.
- Visualize the feature matrix in `visualize_features` by projecting the feature vectors onto the 2D plane. Use Principal Component Analysis for the projection (either use your implementation from assignment 2 or the Julia Package “MultivariateStats”). Plot the projected points and color the features belonging to bicycles in green and those belonging to airplanes in blue.

Discussion

Set $K=5$ for the k -means clustering, and look at the visualized training features. In `answersproblem1.txt`, briefly explain one major problem when setting such a low value for the number of codewords in this classification task.

Submission: Please include only `problem1.jl` and `answersproblem1.txt` in your submission.

Problem 2 - A Multilayer Perceptron (25 points)

In this problem, we will use a multilayer perceptron (MLP) for binary classification:




Figure 2: A simple MLP with 2 inputs, 2 hidden layers, and one output.

For a given training set consisting of N pairs of features and labels (\mathbf{x}_i, y_i) with $y_i \in \{0, 1\}$, we can write the loss function of a MLP as

$$\text{logloss} = -\frac{1}{N} \sum_{i=1}^N (y_i \log(p_i) + (1 - y_i) \log(1 - p_i)), \quad (2)$$

$$p_i = \sigma_2(\mathbf{W}_2 \cdot \sigma_1(\mathbf{W}_1 \cdot \sigma_0(\mathbf{W}_0 \cdot \mathbf{x}_i + \mathbf{b}_0) + \mathbf{b}_1) + \mathbf{b}_2),$$

where \mathbf{W}_k , and \mathbf{b}_k are the weights and biases on the k -th layer respectively. These are the parameters that we want to learn from data. The inputs to every layer undergo a linear transformation and result is subsequently passed through a nonlinear activation function σ_k . Here, we want to use sigmoid activations on all layers, *i.e.*

$$\sigma_k(\mathbf{x}) = \frac{1}{1 + e^{-\mathbf{x}}}. \quad (3)$$

In the following programming assignments we will specify a network architecture by the number of nodes on each layer. For instance the net in Fig. 2 is specified `netdefinition = [2, 4, 4, 1]`. Please solve the following tasks.

Task 1: Training an MLP on toy data

On a 2D toy problem, we will first train an MLP with only 4 neurons on a single hidden layer.

- In the folder you will find two training data sets called `separable.jld` and `nonseparable.jld` containing 2D data points as well as their corresponding class labels. Implement the function `loaddata` that loads the features and the labels.
- To gather a little insight about the data structure, implement the function `showbefore` that shows the 2D data points as a scatter plot. Please color points with label 0 as green and points with label 1 as blue and include a legend that indicates points with corresponding class label.
- Implement the function `initWeight` that initializes all weights and biases of the MLP that structure is defined in the array `netdefinition`. The function should sample from $\mathcal{N}(0, \sigma_W)$ for the weights \mathbf{W} and from $\mathcal{N}(0, \sigma_b)$ for the biases \mathbf{b} .
- Implement two helper functions `weightsToTheta` and `thetaToWeights` that will be used for the optimization. The function `weightsToTheta` concatenates weights and biases into a single vector `theta`. The function `thetaToWeights` inverts `weightsToTheta` by decomposing and reshaping weights and biases from the variable `theta` again.
- Implement the sigmoid activation function and its derivative: `sigmoid` and `dsigmoid_dz`.

- Implement the two functions `nnloss` and `nnlossgrad` that evaluate the loss (forward pass) and its gradient w.r.t. all \mathbf{W}_k and \mathbf{b}_k (backward pass). The gradients can be computed compactly by applying the chain rule for differentiation. Deriving the gradient is a common source of subtle errors. Therefore, we suggest that you first write down the gradient on paper and evaluate it on a simple example. Then compare these numbers to the result of your implementation of the gradient.
- Now train an MLP from the training data by implementing the function `train`. After initializing the weights and biases, you need to minimize the log loss function in eq. (2) w.r.t. its parameters. You can use the function `optimize` of the package “Optim.jl” that implements generic optimization algorithms. You need to pass the loss function, the gradient function, as well as an initial value to `optimize`. Using LBFGS as optimizer should deliver good results.
- Implement the function `predict` that takes a set of data points \mathbf{x} , weight vectors \mathbf{W}_k and offsets \mathbf{b}_k and calculates the network output \mathbf{p} as well as a class label for each data point. If the output of a data point is greater than 0.5, you should assign the class 1 to the i^{th} data point, and class 0 otherwise.
- Write a function `showafter` that plots the data points colored according to their class labels. Additionally, this function should also plot the decision boundary of the trained network. The function `contour` from the PyPlot package and the function `meshgrid` from `Common.jl` might be useful.

Hints: To avoid redundant code you might find it helpful to introduce a separate function for computing the forward pass and keeping track of all necessary quantities like intermediate linear and non-linear activations.

Task 2: Categorize images with an MLP

Now we will train an MLP to classify images of airplanes and motorbikes. If you have not successfully solved `problem1`, do not worry! We have provided you with precomputed features for the training and test images as well as corresponding class labels. The data is provided in the files `imgstrain.jld2` and `imgstest.jld2`. However, note that these do not necessarily represent a valid solution to `problem1`.

Your task is to find a suitable MLP structure with the goal to optimize performance on the test set, *i.e.* achieving a low test error. Try i) different numbers of layers and ii) different numbers of perceptrons in each layer and keep the best value in your solution. What do you observe when changing i) and ii)? Please provide an analysis in `answersproblem2.txt`.

Submission: Please include only `problem2.jl` and `answersproblem2.txt` in your submission.